

Assignment 04

In this assignment, you use Neural Networks with the insurance practice project described at

<https://www.scriptedin.com/contests/view/6>

The data sets can be downloaded, and are provided on HuskyCT.

The computeCost function takes the data set X, for each example/observation in X, does the following (Here L = 3 indicating 3 layers):

I. First the labels are converted:

“1”, which is “No” is converted to a vector of [1,0]; and “2”, which is “Yes” is converted to [0,1] using the identity matrix trick. The new labels are stored in the variable Y.

II. Forward propagation: For each example/observation x


1. $a^{(1)} = x$
2. Add bias value $a^{(1)}_0 = 1$
3. For $l=2 \dots L$
 - a. $z^{(l)} = \theta^{(l-1)} a^{(l-1)}$
 - b. $a^{(l)} = g(z^{(l)})$
 - c. Add bias value $a^{(l)}_0 = 1$
4. End
5. $h_{\theta}(x) = a^{(L)}$

III. Now compute the regularized cost function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K \left(-y^{(i)}_k \log h_{\theta}(x^{(i)})_k - (1 - y^{(i)}_k) \log (1 - h_{\theta}(x^{(i)})_k) \right) + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\theta^{(l)}_{ij})^2$$

IV. Now the function starts back-propagating

The error $\delta_k^{(L)}$ at node k in the output layer L is the difference between the output value $a_k^{(L)}$ and the actual value y_k :



The error made by node j in the layer just before the output layer is

$$\delta_j^{(L-1)} = h'(z_j^{(L-1)}) \sum_{k=1}^K \theta_{kj} \delta_k^{(L)}$$

The sum above is the weighted sum of error for all output nodes, weighted by the weights between the output nodes and node j.

In the case of a tanh function, at layer l, the output:

$$a^{(l)}_j = h(z^{(l)}_j) = \tanh(z^{(l)}_j) = \frac{e^{z^{(l)}_j} - e^{-z^{(l)}_j}}{e^{z^{(l)}_j} + e^{-z^{(l)}_j}}$$

Can prove that:

$$h(z^{(l)}_j) = 1 - \tanh(z^{(l)}_j)^2 = 1 - a^{(l)}_j{}^2$$

so the error above becomes

$$\delta_j^{(L-1)} = (1 - a^{(L-1)}_j{}^2) \sum_{k=1}^{N_k} \theta_{kj} \delta_k^{(L)}$$

N_k is the total number of output nodes at the output layer L

In the case of a logistic function, at layer l, the output:

$$a^{(l)}_j = h(z^{(l)}_j) = \text{sigmoid}(z^{(l)}_j) = \frac{1}{1 + e^{-z^{(l)}_j}}$$

Can prove that:

$$h(z^{(l)}_j) = \text{sigmoid}(z^{(l)}_j) (1 - \text{sigmoid}(z^{(l)}_j)) = a^{(l)}_j (1 - a^{(l)}_j)$$

so the error above becomes

$$\delta_j^{(L-1)} = a^{(l)}_j (1 - a^{(L-1)}_j) \sum_{k=1}^{N_k} \theta_{kj} \delta_k^{(L)}$$

In general, going backward from layer $l+1$ to layer l , the error at node i in layer l is

$$\delta_i^{(l)} = a_i^{(l)}(1 - a_i^{(l)}) \sum_{j=1}^{N_j} \theta_{ji} \delta_j^{(l+1)}$$

It may be better to write the error in the vector form. The error $\delta^{(L)}$ for all output nodes, and their actual values

$$\delta^{(L)} = a^{(L)} - y$$

Vector $\delta^{(L)}$ has all the error $\delta_k^{(L)}$ for layer L in it, and similarly $a^{(L)}$ has all the error $a_k^{(L)}$ for layer L in it, y has all the true values of 0/1 for layer L in it. For other layers, from layer l to layer going backward, the error is

$$\delta^{(l)} = a^{(l)}(1 - a^{(l)}) * \theta^{(l)T} \delta^{(l+1)}$$

The “.” operation is element-by-element multiplication.

The algorithm pseudo code for computing errors at different layers would be, starting from the output layer L

1. $\delta^{(L)} = a^{(L)} - y$
2. For $l=L-1 \dots 2$
 - a. $\delta^{(l)} = a^{(l)}(1 - a^{(l)}) * \theta^{(l)T} \delta^{(l+1)}$
3. End

V. The gradient of the cost function for the instance n is

$$\frac{\partial J_n}{\partial \theta_{ji}} = \frac{\partial J_n}{\partial z_j} \frac{\partial z_j}{\partial \theta_{ji}}$$

But

$$\frac{\partial z_j}{\partial \theta_{ji}} = a_i \text{ (as } z_j = \sum_{i=1}^{N_i} a_i \theta_{ji}, \text{ weighted sum of all the outputs from the previous layer } i \text{)}$$

And

$$\frac{\partial J_n}{\partial z_j} = \delta_j$$

(The layer notation, can be any layer l , is dropped in this equation to make it less confusing), so

$$\frac{\partial J_n}{\partial \theta_{ji}} = \delta_j a_i$$

To compute the gradient of the cost function:

Repeat until convergence:

1. $\Delta^{(l)} = \text{zeros}(\text{shape of } \theta^{(l)}), \text{ for each layer } l=1..L-1$
2. For each instance in the training dataset
 - a. Forward propagation to compute $a^{(l)}$, for each layer $l=1..L$
 - b. Backpropagation to compute $\delta^{(l)}$, for each layer $l=L..2$
 - c. Compute $\Delta^{(l)} = \Delta^{(l)} + \delta^{(l+1)} a^{(l)T}$, for each layer $l=1..L-1$
3. End
4. Compute the regularized gradient $\frac{\partial J(\theta)}{\partial \theta^{(l)}} = \frac{1}{m} \Delta^{(l)} + \lambda \theta^{(l)}$, for each layer $l=1..L-1$

The shape/size of $\Delta^{(l)}$ is the same as the shape/size of $\theta^{(l)}$, which is a matrix. The for loop in step 2 updates the derivatives in the matrix $\Delta^{(l)}$.

The cost and the gradient of the cost are stored in the variables J and grad.

For each layer:

$A^{(l)}$: contains multiple a on its rows, each is corresponding to an example/observation. Note that the last $A^{(L)}$ is also the output H .

$Z^{(l)}$: contains multiple z on its rows, each is corresponding to an example/observation.

delta $\delta^{(l)}$ contains the error on each node at each layer l . BigDelta $\Delta^{(l)}$ contains gradient at layer l .

1. (1 point) Write the `normalize()` function that standardizes/normalizes variables in X
2. (1 point) Write a `sigmoid()` function of z to return a sigmoid/tanh of z
3. (1.5 point) Write a `gradient()` function to return gradient of the nonlinearity
4. (1.5 point) Write the `computeCost()` function that returns the cost
5. (1.5 point) Write the `computeGrad()` function that returns the gradient
6. (1.5 point) Write a `predict()` function to provide prediction.
7. (2 point) Write a `optimize()` function that implements the gradient descent updates to get the optimal θ

The main code was provided, that calls functions above to train a model using the training data and return θ , then test the model using the test data, compute accuracy, confusion matrix, precision, and recall.

Note that in `computeCost()` and `computeGrad()` functions, Theta variables are flattened in `theta` to pass to these functions, then reshaped to the original later in the functions. In the main code, it looks like:

```
theta = np.concatenate(  
    (Theta1.reshape(hidden_num * (1 + input_num), 1, order="F"),  
    Theta2.reshape(label_num * (1 + hidden_num), 1, order="F")))
```

Where `Theta1`, `Theta1` are the weight between the input layer and the hidden layer, and between the hidden layer and the output layer, respectively.

Submission on Husky includes a zip file of:

- a. A notebook named as "group_xxx_assignment_yyy.ipynb"
- b. A Word document/article with detailed explanation
- c. All the other files

Also share the notebook and the article (only after the deadline) to the project site.